

2021 Fall Parallel Functional Programming

Parallel 15 Puzzle Problem

Alex Kuan-yao Huang

kh3120@columbia.edu

Aditya Sidharta

aks2266@columbia.edu

Problem Description

Input: K puzzles with size 2x2 to 4x4

Output: The step used to get the target configuration

Rules:

1. Can only swap the empty cell with the neighbors (up, down, right, left)
2. Not all puzzle is solvable, if not solvable, return -1
3. No duplicated digits appears in the input



8		6
5	4	7
2	3	1

	1	2
3	4	5
6	7	8

Algorithms: this is an **NP hard** problem!

A* Algorithm: heuristic “cost function”, a more balanced algorithm between path length and search space complexity

1. Manhattan Distance
2. Hamming Distance

Breadth-first search: The search space is too large


1. 15 puzzle: $16!/2 = 20922789888000$
2. 24 puzzle: $24!/2 = 7.76 \cdot 10^{24}$

Greedy Algorithm (layer by layer):

Usually not optimal

Optimality: shortest path

Runtime Performance



Algorithms: this is an **NP hard** problem!

A* Algorithm: heuristic “cost function”, a more balanced algorithm between path length and search space complexity

1. Manhattan Distance
2. Hamming Distance

Breadth-first search: The search space is too large


1. 15 puzzle: $16!/2 = 20922789888000$
2. 24 puzzle: $24!/2 = 7.76 \cdot 10^{24}$

Greedy Algorithm (layer by layer):

Usually not optimal

Optimality: shortest path

Runtime Performance



Algorithms: this is an **NP hard** problem!

A* Algorithm: heuristic “cost function”, a more balanced algorithm between path length and search space complexity

1. Manhattan Distance
2. Hamming Distance

8		6
5	4	7
2	3	1

	1	2
3	4	5
6	7	8

Breadth-first search: The search space is too large


1. 15 puzzle: $16!/2 = 20922789888000$
2. 24 puzzle: $24!/2 = 7.76 \cdot 10^{24}$

Greedy Algorithm (layer by layer):

Usually not optimal

Optimality: shortest path

Runtime Performance



Solvability of 15 puzzle problem

Input : nxn puzzle

Output: solvable or not (bool)

Algorithm Solvability :: state -> bool

 If n is odd then

 If number of inversion is even then

 Return true

 Else

 Return false

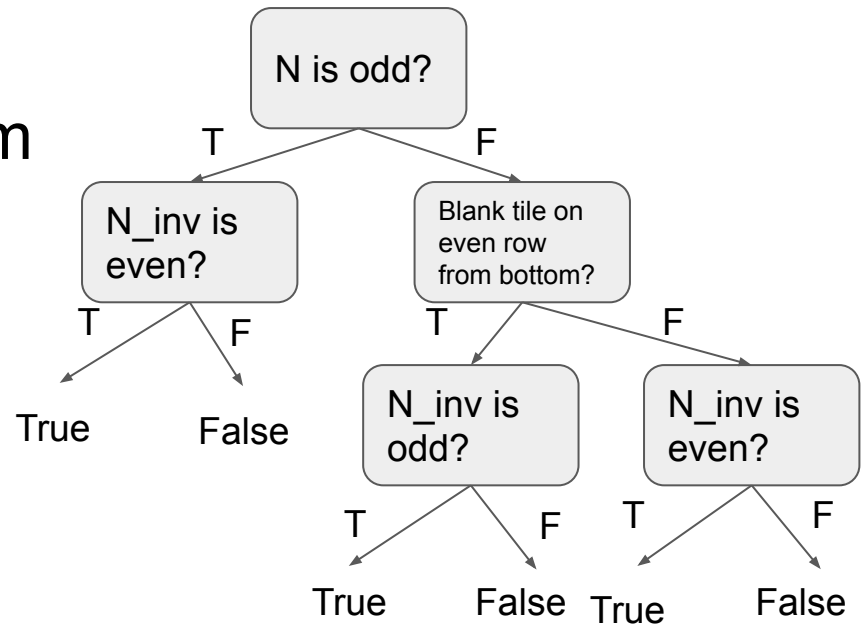
 Else if blank tile is on the even row and number of inversion is odd

 Return true

 Else if blank tile is on the odd row and number of inversion is even

 Return true

 Return false



A* Algorithm

Input : Initial State x K

Output: path length

Algorithm A* :: initialstate -> endstate -> Int

HashMap mp \\ storing visited states

PriorityQueue pq(initialstate, cost, length=0) \\ candidates

While (!pq.empty()) do

 If pq.top().state == endstate:

 Return pq.top().length

 Let neighbors <- getNeighbors pq.top()

 Let validNeighbors <- filter neighbors by mp

 pq.pop()

 For neighbor in validNeighbors do

 Add (neighbor, cost of neighbor, length + 1) to pq

 Add neighbor state to HashMap

Return -1

A* Algorithm

Input : Initial State x K

1. Parallelize solving each puzzle using parList rseq

Output: path length

Algorithm A* :: initialstate -> endstate -> Int

HashMap mp \\ storing visited states

PriorityQueue pq(initialstate, cost, length=0) \\ candidates

While (!pq.empty()) do

 If pq.top().state == endstate:

 Return pq.top().length

 Let neighbors <- getNeighbors pq.top()

 Let validNeighbors <- filter neighbors by mp

 pq.pop()

 For neighbor in validNeighbors do

 Add (neighbor, cost of neighbor, length + 1) to pq

 Add neighbor state to HashMap

Return -1

A* Algorithm

Input : Initial State x K

Output: path length

Algorithm A* :: initialstate -> endstate -> Int

HashMap mp \\ storing visited states

PriorityQueue pq(initialstate, cost, length=0) \\ candidates

While (!pq.empty()) do

 If pq.top().state == endstate:

 Return pq.top().length 2. Parallelize getting four neighbors (will create 4 new states)

 Let neighbors <- getNeighbors pq.top()

 Let validNeighbors <- filter neighbors by mp

 pq.pop()

 For neighbor in validNeighbors do

 Add (neighbor, cost of neighbor, length + 1) to pq

 Add neighbor state to HashMap

Return -1

hashmap1

hashmap2

hashmap3

hashmap4

A* Algorithm

Input : Initial State x K

Output: path length

Algorithm A* :: initialstate -> endstate -> Int

HashMap mp

\\ storing visited states

PriorityQueue pq(initialstate, cost, length=0) \\ candidates

While (!pq.empty()) do

 If pq.top().state == endstate:

 Return pq.top().length

 Let neighbors <- getNeighbors pq.top()

 Let validNeighbors <- filter neighbors by mp





 pq.pop()

 For neighbor in validNeighbors do

 Add (neighbor, cost of neighbor, length + 1) to pq

 Add neighbor state to HashMap

Return -1



3. Creating k priority queues and finding their neighbors in Parallel. Then collect their neighbors' state to hashmap.

Intuition: the optimal route may not be heuristically the best.

Strategy 1: parallelism between test cases

- Similar to Sudoku solution explained in class, it is intuitive to parallelize the solver over different puzzles.
- `parBuffer` was used in order to regulate the number of sparks created to avoid sparks overflow

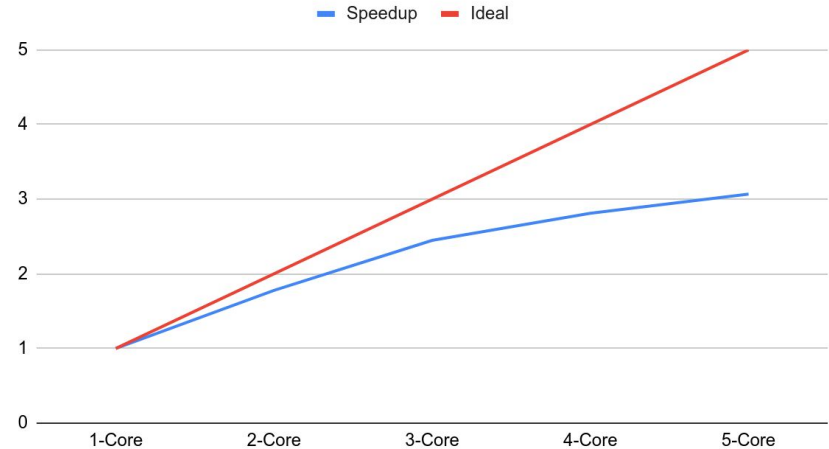
Strategy 1: Parallelism Between Test Cases

We use 100 4x4 solvable puzzles to test our performance on Linux Machine, 6 CPU, 32GB memory

Using parList rseq



	ParallelPuzzle	Speedup
1-Core	12.73	1.00
2-Core	7.16	1.78
3-Core	5.2	2.45
4-Core	4.53	2.81
5-Core	4.15	3.07

Speedup and Ideal



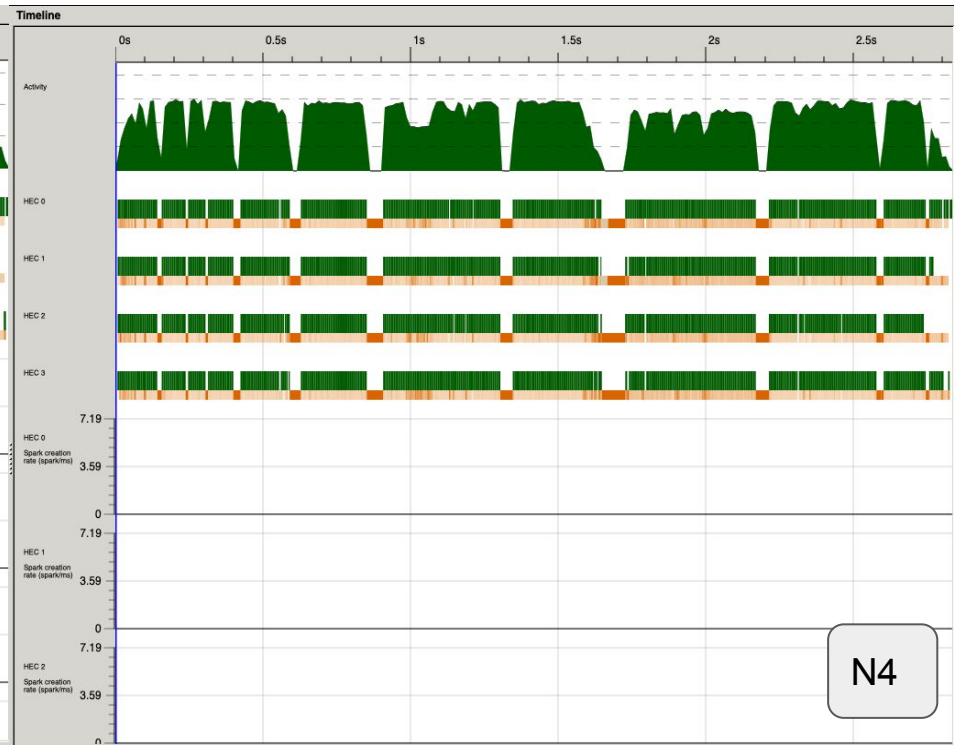

Strategy 1: Parallelism Between Test Cases

The workload is nearly evenly distributed. → Good !



Strategy 1: Parallelism Between Test Cases

The garbage collection time is growing when number of thread is increasing.



Strategy 2: Parallelizing Neighbor Generation

- Most of the steps in A* algorithm depends on previous step
- The only obvious parallelization is the calculation of possible neighbour in each steps


```
getAllNeighborPar :: PuzzleState -> Int -> [PuzzleState]
getAllNeighborPar p n = catMaybes (runEval $ do
  a <- rpar (getUpNeighbor p n)
  b <- rpar (getDownNeighbor p n)
  c <- rpar (getLeftNeighbor p n)
  d <- rpar (getRightNeighbor p n)
  return [a, b, c, d])
```

Strategy 2: Parallelizing Neighbor Generation

	Sequential	ParallelNeighbor
1-Core	13.07	13.15
2-Core	12.89	13.38
3-Core	13.28	13.8
4-Core	13.8	14.7
5-Core	15.17	15.2

As expected, the algorithm does not work well because the calculation of the manhattan distance of a small array (4x4) is insignificant to the overhead from thread / spark creation

Strategy 2: Parallelizing Neighbor Generation



Strategy 3: Parallelism using K priority queue

- Intuition : Not all of the “best” candidate in a certain time-step will result in the optimal solution, as it is possible that the n-th best candidate in a certain time step will lead to the optimal solution
- Thus, putting K-best candidates into different priority queue, and take the result of the fastest thread as the final result might be faster.
- Each of the thread will have their own Hashmap to avoid locking problems.

Strategy 3: Parallelism using K priority queue

Algorithm

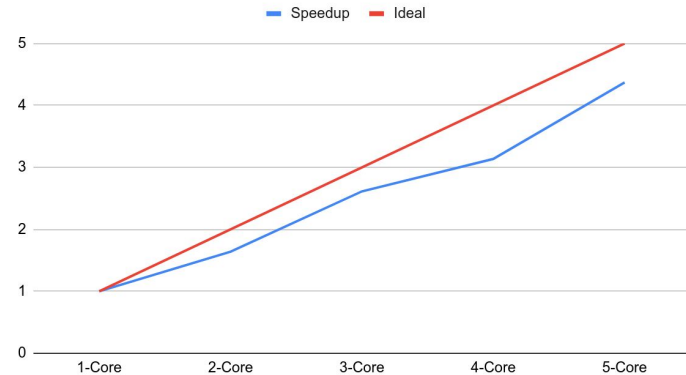
- While the valid neighbours in pq is less than K, run the sequential algorithm for the puzzle
- For (key, prio) in pq.extractMin():
 - Create copy of the current Hashmap hm
 - Create a thread to solve the puzzle, with pq(key, prio) and hm
- Once any of the thread returns the result, kill all other threads and output the result from the finished thread
- Else, if all thread returns non-solvable, return non-solvable

Strategy 3: Parallelism using K priority queue


	Sequential	ParallelPSQ (k=5)
1-Core	13.07	27.54
2-Core	12.89	16.8
3-Core	13.28	10.54
4-Core	13.8	8.78
5-Core	15.17	6.3

	ParallelPSQ (k=5)	Speedup
1-Core	27.54	1.00
2-Core	16.8	1.64
3-Core	10.54	2.61
4-Core	8.78	3.14
5-Core	6.3	4.37

Speedup and Ideal



Strategy 3: Parallelism using K priority queue



Summary

Runtime (s)	Sequential	ParallelNeighbor	ParallelPSQ (k=5)	ParallelPuzzle
1-Core	13.07	13.15	27.54	12.73
2-Core	12.89	13.38	16.8	7.16
3-Core	13.28	13.8	10.54	5.2
4-Core	13.8	14.7	8.78	4.53
5-Core	15.17	15.2	6.3	4.15

Acceleration (compared to sequential)	Sequential	ParallelNeighbor	ParallelPSQ (k=5)	ParallelPuzzle
1-Core	1.00	0.99	0.47	1.03
2-Core	1.01	0.98	0.78	1.83
3-Core	0.98	0.95	1.24	2.51
4-Core	0.95	0.89	1.49	2.89
5-Core	0.86	0.86	2.07	3.15

Future Works

- Implementing Shared memory between the threads to avoid extra computation (Concurrent Hash Map)
- Tune the number of cores and the number of priority queues to get the best result